

Knowledge Distillation Using Early Exit LLMs

Experiments on Confidence Scoring

Edvin 24V0074

Sagar 24D0367

Kapil 210100079

CS 769

Optimization in Machine Learning

20th March

1. Knowledge Distillation

2. Early Exit LLMs

[1] Motivation

References



G. Hinton, O. Vinyals, and J. Dean, 'Distilling the Knowledge in a Neural Network', Mar. 09, 2015, arXiv: arXiv:1503.02531. Accessed: Mar. 06, 2025.

Thank You