

Knowledge Distillation Using Early Exit LLMs

Experiments on Confidence Scoring

Edvin 24V0074

Sagar 24D0367

Kapil 210100079

CS 769

Optimization in Machine Learning

20 March 2016

1. Knowledge Distillation

2. Early Exit LLMs

[Hinton et al., 2015] Motivation

References



Hinton, G., Vinyals, O., and Dean, J. (2015).
Distilling the Knowledge in a Neural Network.
arXiv:1503.02531 [cs, stat].

Thank You